

# Braiding the Blockchain


Bob McElrath, Ph.D.

Scaling Bitcoin  
Hong Kong  
Dec 7 2015  
[bmcelrath@sldx.com](mailto:bmcelrath@sldx.com)



# Save the Orphans!


Orphans are *not* a necessary component of the operation of bitcoin!



- Orphans occur when miners do not know about the existence of another block (B) before generating theirs (C)
- Simultaneous block generation is *unavoidable*
- It doesn't require us to deprive a miner of profit

# Save the Orphans!


Orphans are *not* a necessary component of the operation of bitcoin!



- What if block C contains no conflicting transactions?
- What if block C contains a duplicate transaction?
- *There is no conflict so let us call C a sibling.*

# Save the Orphans!


Orphans are *not* a necessary component of the operation of bitcoin!



- A future block must be able to tie up B and C, indicating that there is no conflict.
- To get rid of orphans, *blocks must have multiple parents*

# Save the Orphans!

Orphans are *not* a necessary component of the operation of bitcoin!



- If C contains a double-spend relative to B, then C forms a new chain tip.

# Down the Rabbit Hole

- Allowing blocks to have multiple parents creates a data structure called a *Directed Acyclic Graph*.
- What if I just throw out blocks as fast as I desire, do algorithms exist that could make sense of the chaos and define a highest work “tip”?
- The blockchain is an over-simplified data structure, with some unfortunate consequences (orphans, selfish mining).



# Down the Rabbit Hole

- Allowing blocks to have multiple parents creates a data structure called a *Directed Acyclic Graph*.
- What if I just throw out blocks as fast as I desire, do algorithms exist that could make sense of the chaos and define a highest work “tip”?
- The blockchain is an over-simplified data structure, with some unfortunate consequences (orphans, selfish mining).



# Down the Rabbit Hole

- Allowing blocks to have multiple parents creates a data structure called a *Directed Acyclic Graph*.
- What if I just throw out blocks as fast as I desire, do algorithms exist that could make sense of the chaos and define a highest work “tip”?
- The blockchain is an over-simplified data structure, with some unfortunate consequences (orphans, selfish mining).





# The Directed Acyclic Graph

Allowing blocks to have multiple parents creates a:


**Directed** Blocks have parents, parents cannot refer to children

**Acyclic** A cycle is cryptographically impossible

**Graph** Structure is non-linear (no “height”)

- A DAG can be *partial ordered* in linear time.
- We have to make a restriction relative to a more general dag, so I'm going to name this data structure a *braid*.

## Braid Terminology



**Braid** A Directed Acyclic Graph having no *incest* (no triangles)

**Bead** Analog of Bitcoin's blocks (green circles)

**Sibling** A *bead* that cannot be partial ordered relative to myself:  
the pairs (B,C) and (E,F)

**Incest** A parent that is simultaneously an ancestor of another parent (**disallowed**)

# Our Approach

- Because of *selfish mining*<sup>1</sup> we will incentivize miners to quickly transmit beads
- Because GHOST<sup>2</sup> worsens some attacks, we will require that *parents must not contain conflicting transactions*
- Allow all of these to be decided per-node:
  - Bead time; Bead target difficulty; Bead size
- Assume Braids will be a parallel, faster layer to Bitcoin blocks:
  - Beads will be constructed such that they are valid Bitcoin blocks, if they meet bitcoin's difficulty target.
- Publish beads *ex-post-facto* because knowing who the current leader is (Bitcoin-NG<sup>3</sup>) opens a new vulnerability


---

<sup>1</sup>Eyal, Sirer, arXiv:1311.0243

<sup>2</sup>Sompolinsky, Zohar, [ia.cr/2013/881](http://ia.cr/2013/881)

<sup>3</sup>Eyal, Gencer, Sirer, Renesse, arXiv:1510.02037

# Braid Fork Example



- A double-spend occurs in A and B  
⇒ We must evaluate which braid has the most work
- Beads in each fork reference either A or B as a parent
- The highest-work braid will be decided by evaluating the work in the combined work of all beads in the red and blue subgraphs.

# How to Incentivize Miners

Miner incentives must be aligned with correct operation of the network

- Consensus is *created* by the profit-maximizing miners
- Let the *reward* be proportional to a miner's target difficulty
  - ⇒ Miners can individually choose target and block rate based on *other considerations* (e.g. bandwidth)
  - ⇒ Bandwidth and CPU is now the only limiting factor for the network!
- The existence of siblings/orphans means *we cannot decide miner coin allocation until all beads are seen by all nodes*
  - ⇒ Coin allocation will be calculated 100 blocks later

We will define several quantities that we will use in a new miner incentive formula (a.k.a. How many bitcoins do I get?)

# Siblings

A sibling  $S$  is an analog to Bitcoin *orphans*. It is defined as

A bead that cannot be ordered to come before or after mine using only the DAGs partial order

- Siblings are defined *per braid tip*
- Siblings *must not* contain conflicting transactions
- Siblings *may* contain duplicate transactions

If siblings share the same transaction, each sibling will be allocated a work-weighted fraction of the tx fee. (e.g. 2 siblings at the same target difficulty will each receive 1/2 of the tx fee)


# Cohort Difficulty

The *cohort difficulty*  $D$  is the work of other miners during the time window in which I was mining. It is defined as:

The combined work of all beads between  
my youngest parent and my oldest child.

- A miner with large cohort difficulty relative to his own is playing games or following a perverse incentive
    - Trying to steal fees by becoming everyone's sibling
    - Withholding blocks (children are late)
- ⇒ Incentivizing small cohort difficulty incentivizes fast block transmission

# Cohorts and Siblings



The cohort of bead B is (D,E,G) while its siblings are (D,E)

Quiz:

- Siblings of G? Cohort of G?
- Siblings of D? Cohort of D?



# Miner Incentive Formula

The miner of block  $i$  receives a reward  $R_i$ :

$$R_i = \sum_t^{T_i} f_t \frac{d_i}{D_i} \left( \frac{1}{N_t} \right) + C \frac{d_i}{D_i} \left( \frac{1}{N} \right); \quad N_t = \sum_s^{S_t} \frac{d_s}{D_s}; \quad N = \sum_j^{N_c} \frac{d_j}{D_j}$$

$i, j$  bead indices

$t$  transaction index

$d_i$  difficulty = 1/target

$D_i$  the cohort difficulty

$T_i$  number of transactions in bead  $i$

$C$  block reward = 25 BTC

$f_t$  Transaction fee for tx  $t$

$S_t$  number of siblings containing tx  $t$  ( $S_t = 1$  if no siblings)

$N_t$  Sum of weighted difficulty over siblings containing tx  $t$

$N$  Sum of weighted difficulty over all beads (normalization)

# Miner Incentive Formula

The miner of block  $i$  receives a reward  $R_i$ :

$$R_i = \sum_t^{T_i} f_t \frac{d_i}{D_i} \left( \frac{1}{N_t} \right) + C \frac{d_i}{D_i} \left( \frac{1}{N} \right); \quad N_t = \sum_s^{S_t} \frac{d_s}{D_s}; \quad N = \sum_j^{N_c} \frac{d_j}{D_j}$$

This miner incentive formula is constructed such that it is:

- *linear* in miner difficulty  $d_i$  (miners set their own target)
  - For  $\frac{d_i}{D_i} \ll N$  miner income is independent of target
  - Smaller  $d_i$  means *smoother* income distribution over time
- Fair (difficulty-weighted) split of fees  $f_t$  among siblings

# Miner Incentive Formula

The miner of block  $i$  receives a reward  $R_i$ :

$$R_i = \sum_t^{T_i} f_t \frac{d_i}{D_i} \left( \frac{1}{N_t} \right) + C \frac{d_i}{D_i} \left( \frac{1}{N} \right); \quad N_t = \sum_s^{S_t} \frac{d_s}{D_s}; \quad N = \sum_j^{N_c} \frac{d_j}{D_j}$$

Consequences of this incentive structure:

- We're incentivized to *optimize the p2p topology* to quickly propagate blocks
- Use of cohort difficulty  $D_i$  incentivizes fast transmission of blocks
- Small miners can mine without joining pools: coinbase has many outputs (like p2pool)

# Evaluating the Best Braid

The sub-braid containing the most work can be determined by estimating the hash rate

$$H = \sum_i^{\text{miners}} H_i$$

The best way to do this is using a likelihood function:

$$W_{\alpha}(\{H_i\}) = -\log L = -\log \prod_i^{\{x_i\}} P_{x_i}(H_i t, k_i)$$

where  $P_x(h, k)$  is the Poisson distribution

- Miners are incentivized to include all chain tips as parents because it gives the sub-dag they're mining on more work.

# Confirmation Times

How do I know when a transaction is “confirmed”?

- Satoshi’s analysis still applies, and we must keep Bitcoin’s payout schedule
  - ⇒ Counting six bitcoin blocks is still resonable
- Much better analyses are possible
  - Do there exist other braid-tips with similar work?
  - What’s the ratio of work in my braid tip and the next closest?
  - Has the hash power recently changed?

*We have much more data: I’d like to see a whole class of risk evaluation methods added for different users and use cases.*

# Confirmation Times

How do I know when a transaction is “confirmed”?

- Satoshi’s analysis still applies, and we must keep Bitcoin’s payout schedule
  - ⇒ Counting six bitcoin blocks is still resonable
- Much better analyses are possible
  - Do there exist other braid-tips with similar work?
  - What’s the ratio of work in my braid tip and the next closest?
  - Has the hash power recently changed?

*We have much more data: I’d like to see a whole class of risk evaluation methods added for different users and use cases.*

# Conclusions

- Getting rid of orphans forces us to the Braid structure.
- Transaction volume is limited only by bandwidth and CPU!
- Confirmation times can be *much* faster, limited only by the propagation time to reach the entire network (e.g. the size of the Earth).
- Algorithms are more complex, but seem to all be  $O(N)$ : We don't have to solve the Travelling Salesman Problem.
- Miner income becomes much smoother and more predictable
- Many ways to put this into bitcoin many incentive models: simulation and testing is necessary.
- Smaller miners do not have to use pools
  - ⇒ mining decentralization!